


From: Madhushree S Shetty madhushree.cs22@bmsce.ac.in 
Subject: Lab08
Date: 6 May 2025 at 3:45 PM
To: Madhushree S Shetty madhushreesetty2003@gmail.com

MS

```
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/eclipse-workspace/Lab08

hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/Desktop
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/Desktop
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/Desktop

hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/Desktop$ cd /home/hadoop/eclipse-workspace/Weather/src
bash: cd: /home/hadoop/eclipse-workspace/Weather/src: No such file or directory
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/Desktop$ hadoop fs -mkdir -p /user/hadoop/input
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/Desktop$ hadoop fs -copyFromLocal -f /home/hadoop/Desktop/1901 /user/hadoop/input/data.txt
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/Desktop$ hadoop fs -ls /user/hadoop/input
Found 3 items
-rw-r--r-- 1 hadoop supergroup      888190 2025-05-06 15:11 /user/hadoop/input/1901
-rw-r--r-- 1 hadoop supergroup      888190 2025-05-06 15:36 /user/hadoop/input/data.txt
-rw-r--r-- 1 hadoop supergroup       90 2025-05-06 14:45 /user/hadoop/input/sample.txt
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/Desktop$ hadoop fs -rm -r /user/hadoop/output
Deleted /user/hadoop/output
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/Desktop$ cd /home/hadoop/eclipse-workspace/Weather
bash: cd: /home/hadoop/eclipse-workspace/Weather: No such file or directory
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/Desktop$ AC
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/Desktop$ hadoop jar Weather.jar TempAnalysisDriver /user/hadoop/input /user/hadoop/output
JAR does not exist or is not a normal file: /home/hadoop/Desktop/Weather.jar
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/Desktop$ hadoop jar /home/hadoop/eclipse-workspace/Weather.jar TempAnalysisDriver /user/hadoop/input /user/hadoop/output
2025-05-06 15:37:33,018 INFO impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2025-05-06 15:37:33,054 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2025-05-06 15:37:33,054 INFO impl.MetricsSystemImpl: JobTracker metrics system started
2025-05-06 15:37:33,178 INFO input.FileInputFormat: Total input files to process : 3
2025-05-06 15:37:33,187 INFO mapreduce.JobSubmitter: number of splits:3
2025-05-06 15:37:33,249 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local774182108_0001
2025-05-06 15:37:33,249 INFO mapreduce.JobSubmitter: Executing with tokens: []
2025-05-06 15:37:33,307 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
2025-05-06 15:37:33,307 INFO mapreduce.Job: Running job: job_local774182108_0001
2025-05-06 15:37:33,308 INFO mapred.LocalJobRunner: OutputCommitter set in config null
2025-05-06 15:37:33,311 INFO output.PathOutputCommitterFactory: No output committer factory defined, defaulting to FileOutputCommitterFactory
2025-05-06 15:37:33,312 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2025-05-06 15:37:33,312 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2025-05-06 15:37:33,312 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter
2025-05-06 15:37:33,355 INFO mapred.LocalJobRunner: Waiting for map tasks
2025-05-06 15:37:33,355 INFO mapred.LocalJobRunner: Starting task: attempt_local774182108_0001_m_000000_0
2025-05-06 15:37:33,368 INFO output.PathOutputCommitterFactory: No output committer factory defined, defaulting to FileOutputCommitterFactory
2025-05-06 15:37:33,368 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2025-05-06 15:37:33,368 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2025-05-06 15:37:33,375 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
2025-05-06 15:37:33,378 INFO mapred.MapTask: Processing split: hdfs://localhost:9000/user/hadoop/input/1901:0+888190
2025-05-06 15:37:33,411 INFO mapred.MapTask: (EQUATOR) 0 kvl 26214396(104857584)
2025-05-06 15:37:33,411 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100
2025-05-06 15:37:33,411 INFO mapred.MapTask: soft limit at 83886080
Merged Map outputs:
GC time elapsed (ms)=0
Total committed heap usage (bytes)=1052770304
File Input Format Counters
Bytes Read=1776380
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/Desktop$ hadoop fs -ls /user/hadoop/output
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/Desktop$ cd ~/eclipse-workspace/Weather/src
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/Desktop$ cd ~/eclipse-workspace/Weather/src
hadoop com.sun.tools.javac.Main *.java
jar cf Weather.jar *.class
mv Weather.jar ..
bash: cd: /home/hadoop/eclipse-workspace/Weather/src: No such file or directory
error: file not found: *.java
Usage: javac <options> <source files>
use --help for a list of possible options
*.class : no such file or directory
mv: cannot stat 'Weather.jar': No such file or directory
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/Desktop$ cd ~/eclipse-workspace/Lab08/src
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/eclipse-workspace/Lab08/src$ hadoop com.sun.tools.javac.Main *.java
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/eclipse-workspace/Lab08/src$ jar cf Weather.jar *.class
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/eclipse-workspace/Lab08/src$ mv Weather.jar ..
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/eclipse-workspace/Lab08/src$ hadoop fs -mkdir -p /user/hadoop/aaa
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/eclipse-workspace/Lab08/src$ hadoop fs -copyFromLocal -f /home/hadoop/Desktop/1901 /user/hadoop/aaa/text.txt
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/eclipse-workspace/Lab08/src$ hadoop fs -rm -r /user/hadoop/aaa/output
Deleted /user/hadoop/aaa/output
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/eclipse-workspace/Lab08/src$ cd ~/eclipse-workspace/Lab08
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/eclipse-workspace/Lab08$ hadoop jar Weather.jar TempAnalysisDriver /user/hadoop/aaa/text.txt /user/hadoop/aaa/output
2025-05-06 15:42:33,864 INFO impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2025-05-06 15:42:33,903 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2025-05-06 15:42:33,903 INFO impl.MetricsSystemImpl: JobTracker metrics system started
2025-05-06 15:42:34,007 INFO input.FileInputFormat: Total input files to process : 1
2025-05-06 15:42:34,031 INFO mapreduce.JobSubmitter: number of splits:1
2025-05-06 15:42:34,097 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local1383725974_0001
2025-05-06 15:42:34,098 INFO mapreduce.JobSubmitter: Executing with tokens: []
2025-05-06 15:42:34,152 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
2025-05-06 15:42:34,153 INFO mapreduce.Job: Running job: job_local1383725974_0001
2025-05-06 15:42:34,153 INFO mapred.LocalJobRunner: OutputCommitter set in config null
2025-05-06 15:42:34,158 INFO output.PathOutputCommitterFactory: No output committer factory defined, defaulting to FileOutputCommitterFactory
2025-05-06 15:42:34,158 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2025-05-06 15:42:34,158 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2025-05-06 15:42:34,159 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter
2025-05-06 15:42:34,206 INFO mapred.LocalJobRunner: Waiting for map tasks
2025-05-06 15:42:34,207 INFO mapred.LocalJobRunner: Starting task: attempt_local1383725974_0001_m_000000_0
2025-05-06 15:42:34,220 INFO output.PathOutputCommitterFactory: No output committer factory defined, defaulting to FileOutputCommitterFactory
2025-05-06 15:42:34,220 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2025-05-06 15:42:34,220 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2025-05-06 15:42:34,227 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
2025-05-06 15:42:34,230 INFO mapred.MapTask: Processing split: hdfs://localhost:9000/user/hadoop/aaa/text.txt:0+000100
Merged Map outputs:
GC time elapsed (ms)=0
Total committed heap usage (bytes)=1052770304
File Output Format Counters
Bytes Written=88
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/eclipse-workspace/Lab08$ hadoop fs -cat /user/hadoop/aaa/output/part-r-00000
0000 Avg: 99 Max: 206
0029 Avg: 50 Max: 9999
0035 Avg: 15 Max: 267
0048 Avg: 83 Max: 89
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: ~/eclipse-workspace/Lab08$
```